

R Project Airbnb

Anupriya Singh

2024-07-04

R Markdown

```
# Data Importing
listings <- read.csv("D:/r programming/listings.csv")
View(listings)
```

```
## Codes with outputs
```

```
library("dplyr")
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library("tidyr")
```

```
## Warning: package 'tidyr' was built under R version 4.3.3
```

```
library("ggplot2")
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
library("lubridate")
```

```
## Warning: package 'lubridate' was built under R version 4.3.3
```

```
##
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
```

```
colSums(is.na(listings)) # N/A are in -> price, reviews_per_month
```

```
##           id           name
##           0           0
##      host_id      host_name
##           0           0
##   neighbourhood      latitude
##           0           0
##      longitude      room_type
##           0           0
##           price      minimum_nights
##           50           0
##   number_of_reviews      last_review
##           0           0
##   reviews_per_month calculated_host_listings_count
##           45           0
##   availability_365      number_of_reviews_ltm
##           0           0
```

```
# Replace NA in reviews_per_month with 0
```

```
listings <- mutate(listings, reviews_per_month = ifelse(is.na(reviews_per_month), 0, reviews_per_month))
```

```
# Converting last_review as date format
```

```
listings <- mutate(listings, last_review = parse_date_time(last_review, orders = c("Ymd", "Y-m-d", "dmy"))
```

```
# Drop rows with missing values in other crucial columns
```

```
listings <- drop_na(listings, price, minimum_nights, room_type, latitude, longitude, last_review)
```

```
# Transforming data
```

```
listings <- mutate(listings, room_type = as.factor(room_type))
```

```
listings <- mutate(listings, reviews_per_year = reviews_per_month * 12)
```

```
# Exploratory Data Analysis
```


```
summary(listings)
```

```
##           id           name           host_id           host_name
##   Min.   :1.489e+06   Length:313   Min.    : 649068   Length:313
##   1st Qu.:4.733e+07   Class :character   1st Qu.: 47625981   Class :character
##   Median :6.686e+17   Mode  :character   Median :172539578   Mode  :character
##   Mean    :5.052e+17           Mean    :215274022
##   3rd Qu.:9.022e+17           3rd Qu.:382970529
##   Max.    :1.147e+18           Max.    :552465537
##   neighbourhood      latitude      longitude      room_type
##   Length:313         Min.    :42.63   Min.    :-73.83   Entire home/apt:231
##   Class :character    1st Qu.:42.65   1st Qu.: -73.79   Private room    : 81
##   Mode  :character    Median :42.66   Median : -73.77   Shared room     : 1
##                       Mean    :42.66   Mean    :-73.78
##                       3rd Qu.:42.66   3rd Qu.: -73.76
##                       Max.    :42.69   Max.    :-73.74
```

```
##      price      minimum_nights  number_of_reviews
## Min.   : 26.0    Min.   : 1.000    Min.   : 1.00
## 1st Qu.: 79.0    1st Qu.: 1.000    1st Qu.: 8.00
## Median : 105.0   Median : 1.000    Median : 29.00
## Mean   : 125.3   Mean   : 3.489    Mean   : 70.32
## 3rd Qu.: 136.0   3rd Qu.: 2.000    3rd Qu.: 84.00
## Max.   :1379.0   Max.   :30.000    Max.   :795.00
## last_review      reviews_per_month
## Min.   :2018-07-27 00:00:00.00    Min.   : 0.040
## 1st Qu.:2024-02-19 00:00:00.00    1st Qu.: 0.690
## Median :2024-04-14 00:00:00.00    Median : 1.720
## Mean   :2024-02-05 05:17:26.65    Mean   : 2.283
## 3rd Qu.:2024-04-26 00:00:00.00    3rd Qu.: 3.160
## Max.   :2024-05-06 00:00:00.00    Max.   :11.050
## calculated_host_listings_count availability_365 number_of_reviews_ltm
## Min.   : 1.00                      Min.   : 3.0    Min.   : 0.00
## 1st Qu.: 1.00                      1st Qu.:126.0   1st Qu.: 4.00
## Median : 3.00                      Median :243.0   Median : 14.00
## Mean   : 5.54                      Mean   :220.4    Mean   : 21.21
## 3rd Qu.: 9.00                      3rd Qu.:318.0   3rd Qu.: 29.00
## Max.   :22.00                      Max.   :365.0    Max.   :131.00
## reviews_per_year
## Min.   : 0.48
## 1st Qu.: 8.28
## Median : 20.64
## Mean   : 27.40
## 3rd Qu.: 37.92
## Max.   :132.60
```

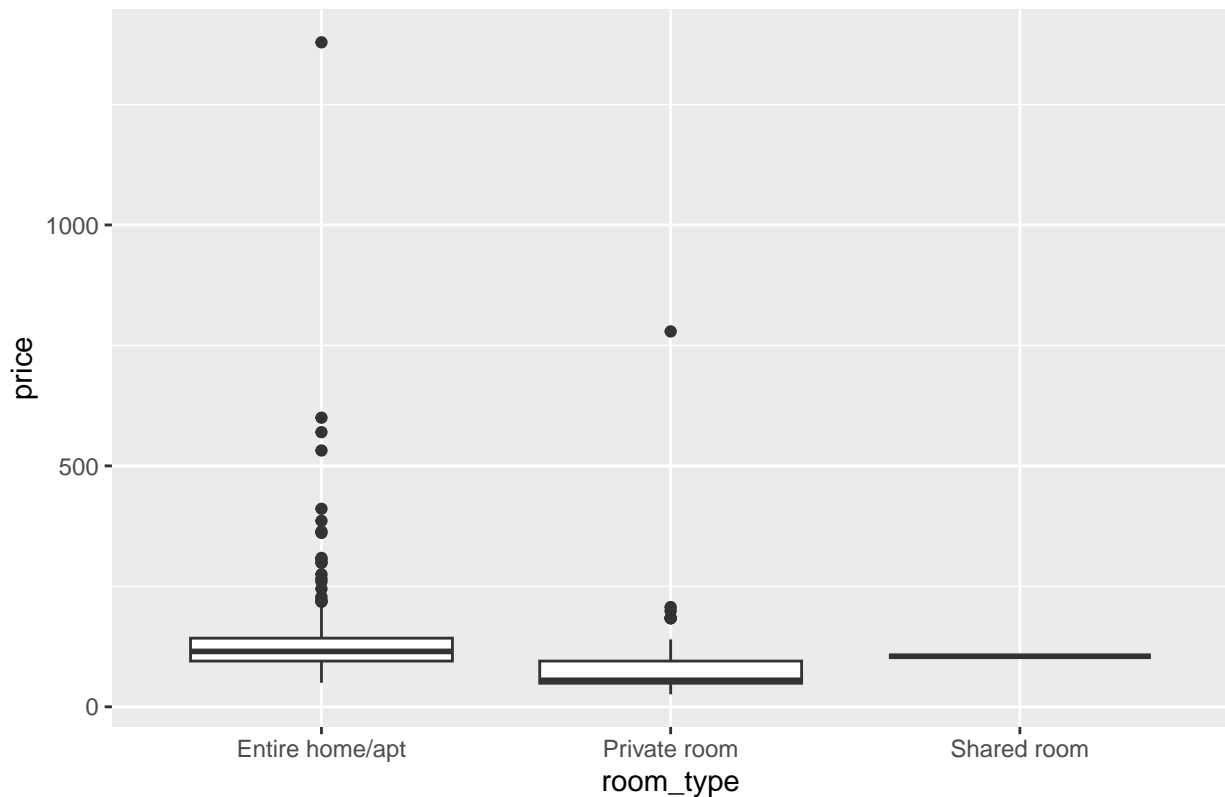
```
# Visualizations
ggplot(listings, aes(x = price)) +
  geom_histogram(binwidth = 50) +
  ggtitle("Distribution of Prices")
```

Distribution of Prices



```
ggplot(listings, aes(x = room_type, y = price)) +  
  geom_boxplot() +  
  ggtitle("Price by Room Type")
```

Price by Room Type




```
# Exploratory Data Analysis
summary(listings)
```

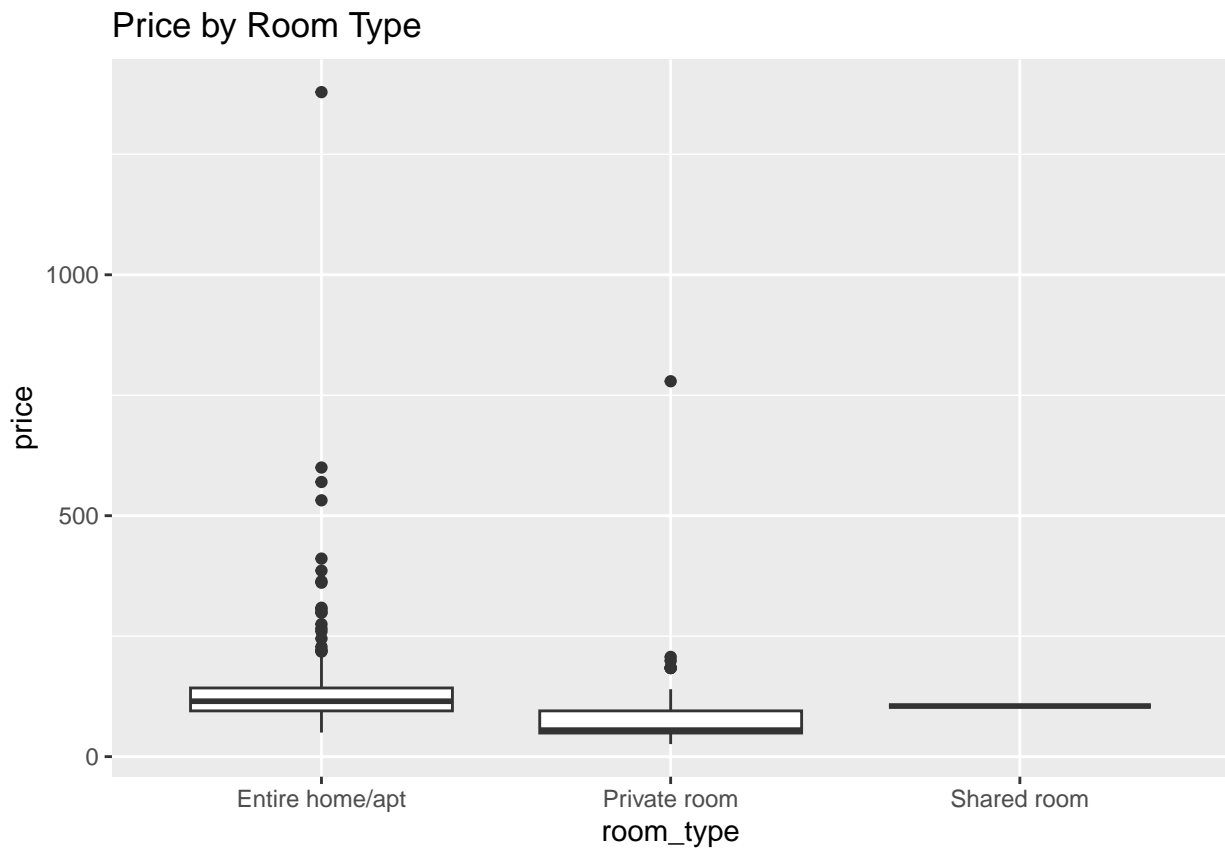
```
##          id          name          host_id          host_name
##  Min.   :1.489e+06  Length:313    Min.    : 649068  Length:313
##  1st Qu.:4.733e+07  Class :character  1st Qu.: 47625981  Class :character
##  Median :6.686e+17  Mode  :character  Median :172539578  Mode  :character
##  Mean   :5.052e+17                                Mean   :215274022
##  3rd Qu.:9.022e+17                                3rd Qu.:382970529
##  Max.   :1.147e+18                                Max.   :552465537
## neighbourhood    latitude    longitude    room_type
## Length:313       Min.    :42.63  Min.    : -73.83  Entire home/apt:231
## Class :character  1st Qu.:42.65  1st Qu.: -73.79  Private room   : 81
## Mode  :character  Median :42.66  Median : -73.77  Shared room    :  1
##                  Mean   :42.66  Mean   : -73.78
##                  3rd Qu.:42.66  3rd Qu.: -73.76
##                  Max.   :42.69  Max.   : -73.74
## price            minimum_nights  number_of_reviews
##  Min.   : 26.0    Min.    : 1.000    Min.    : 1.00
##  1st Qu.: 79.0    1st Qu.: 1.000    1st Qu.: 8.00
##  Median :105.0    Median : 1.000    Median :29.00
##  Mean   :125.3    Mean   : 3.489    Mean   :70.32
##  3rd Qu.:136.0    3rd Qu.: 2.000    3rd Qu.:84.00
##  Max.   :1379.0   Max.   :30.000    Max.   :795.00
## last_review      reviews_per_month
##  Min.   :2018-07-27 00:00:00.00  Min.    : 0.040
```

```
## 1st Qu.:2024-02-19 00:00:00.00 1st Qu.: 0.690
## Median :2024-04-14 00:00:00.00 Median : 1.720
## Mean :2024-02-05 05:17:26.65 Mean : 2.283
## 3rd Qu.:2024-04-26 00:00:00.00 3rd Qu.: 3.160
## Max. :2024-05-06 00:00:00.00 Max. :11.050
## calculated_host_listings_count availability_365 number_of_reviews_ltm
## Min. : 1.00 Min. : 3.0 Min. : 0.00
## 1st Qu.: 1.00 1st Qu.:126.0 1st Qu.: 4.00
## Median : 3.00 Median :243.0 Median : 14.00
## Mean : 5.54 Mean :220.4 Mean : 21.21
## 3rd Qu.: 9.00 3rd Qu.:318.0 3rd Qu.: 29.00
## Max. :22.00 Max. :365.0 Max. :131.00
## reviews_per_year
## Min. : 0.48
## 1st Qu.: 8.28
## Median : 20.64
## Mean : 27.40
## 3rd Qu.: 37.92
## Max. :132.60
```

```
# Visualizations
ggplot(listings, aes(x = price)) +
  geom_histogram(binwidth = 50) +
  ggtitle("Distribution of Prices")
```



```
ggplot(listings, aes(x = room_type, y = price)) +
  geom_boxplot() +
  ggtitle("Price by Room Type")
```



```
# Feature Engineering
# Creating a variable which contains distance of a place from the given landmark
landmark_lat <- 40.748817
landmark_lon <- -73.985428

# Define a function to calculate the Haversine distance
haversine_distance <- function(lat1, lon1, lat2, lon2) {
  # Convert degrees to radians
  radians <- pi / 180
  lat1 <- lat1 * radians
  lon1 <- lon1 * radians
  lat2 <- lat2 * radians
  lon2 <- lon2 * radians

  # Haversine formula
  dlat <- lat2 - lat1
  dlon <- lon2 - lon1
  a <- sin(dlat / 2)^2 + cos(lat1) * cos(lat2) * sin(dlon / 2)^2
  c <- 2 * atan2(sqrt(a), sqrt(1 - a))

  # Radius of Earth in kilometers
```

```

R <- 6371
distance <- R * c
return(distance)
}

listings <- mutate(listings,distance_from_landmark = haversine_distance(latitude, longitude, landmark_1

# Modeling
library(caret)

```

```
## Warning: package 'caret' was built under R version 4.3.3
```

```
## Loading required package: lattice
```

```

set.seed(123)
train_indices <- sample(seq_len(nrow(listings)), size = 0.7 * nrow(listings))
train_data <- listings[train_indices, ]
test_data <- listings[-train_indices, ]

# Exclude 'name' and 'host_name' variables from the model
train_data <- select(train_data, -name, -host_name)
test_data <- select(test_data, -name, -host_name)

# Building a regression model
model <- lm(price ~ ., data = train_data)
summary(model)

```

```

##
## Call:
## lm(formula = price ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -158.14  -41.65   -6.94   22.84  1129.74
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.189e+06  9.229e+06   0.129  0.8976
## id             8.861e-18  2.312e-17   0.383  0.7019
## host_id        1.026e-07  4.432e-08   2.315  0.0217 *
## neighbourhoodeLEVENTH WARD -8.476e+01  8.083e+01  -1.049  0.2957
## neighbourhoodeFIFTEENTH WARD -1.566e+02  8.563e+01  -1.829  0.0689 .
## neighbourhoodeFIFTH WARD -1.298e+02  1.017e+02  -1.276  0.2034
## neighbourhoodeFIRST WARD  2.431e+01  8.731e+01   0.278  0.7810
## neighbourhoodeFOURTEENTH WARD -8.238e+01  6.073e+01  -1.356  0.1766
## neighbourhoodeFOURTH WARD -2.236e+01  1.276e+02  -0.175  0.8611
## neighbourhoodeNINTH WARD -6.421e+01  6.349e+01  -1.011  0.3132
## neighbourhoodeSECOND WARD  1.901e+01  8.668e+01   0.219  0.8266
## neighbourhoodeSEVENTH WARD -2.291e+01  7.278e+01  -0.315  0.7532
## neighbourhoodeSIXTH WARD -3.959e+01  7.530e+01  -0.526  0.5996
## neighbourhoodeTENTH WARD -5.665e+01  6.764e+01  -0.838  0.4033
## neighbourhoodeTHIRD WARD -4.857e+01  9.619e+01  -0.505  0.6142

```



```
## neighbourhoodTHIRTEENTH WARD -1.421e+02 7.394e+01 -1.922 0.0560 .
## neighbourhoodTWELFTH WARD -1.435e+02 9.550e+01 -1.503 0.1345
## latitude -3.694e+04 2.529e+05 -0.146 0.8840
## longitude -4.190e+03 1.463e+04 -0.286 0.7749
## room_typePrivate room -8.152e+01 1.885e+01 -4.324 2.47e-05 ***
## room_typeShared room -3.695e+01 1.168e+02 -0.316 0.7521
## minimum_nights -4.964e-01 1.232e+00 -0.403 0.6875
## number_of_reviews 4.138e-02 1.315e-01 0.315 0.7533
## last_review -5.125e-07 5.104e-07 -1.004 0.3166
## reviews_per_month -9.339e+00 7.881e+00 -1.185 0.2375
## calculated_host_listings_count 2.923e+00 1.413e+00 2.069 0.0399 *
## availability_365 1.066e-01 6.984e-02 1.526 0.1286
## number_of_reviews_ltm 3.528e-01 5.825e-01 0.606 0.5455
## reviews_per_year NA NA NA NA
## distance_from_landmark 3.675e+02 2.286e+03 0.161 0.8724
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 102.8 on 190 degrees of freedom
## Multiple R-squared: 0.2082, Adjusted R-squared: 0.09157
## F-statistic: 1.785 on 28 and 190 DF, p-value: 0.01271
```

```
# 6. Model Evaluation
# Predicting on test set
predictions <- predict(model, newdata = test_data)

# Calculate RMSE
rmse <- sqrt(mean((predictions - test_data$price)^2))
print(paste("RMSE: ", rmse))
```

```
## [1] "RMSE: 101.554272164889"
```

```
# Visualizing model performance
ggplot(data.frame(Predicted = predictions, Actual = test_data$price), aes(x = Predicted, y = Actual)) +
  geom_point() +
  geom_abline(slope = 1, intercept = 0, color = 'blue') +
  ggtitle("Predicted vs Actual Prices")
```

